

# What is tidy data?

RESHAPING DATA WITH TIDYR



**Jeroen Boeye**

Head of Machine Learning, Faktion

*Happy families are all alike, but every unhappy family is unhappy in its own way.*

---

**Leo Tolstoy**

*Tidy datasets are all alike, but every messy dataset is messy in its own way.*

---

**Hadley Wickham**

# Rectangular data

## Structure

- Columns
- Rows
- Cells


| name           | homeworld | species |
|----------------|-----------|---------|
| Luke Skywalker | Tatooine  | Human   |
| R2-D2          | Naboo     | Droid   |
| Darth Vader    | Tatooine  | Human   |
| Obi-Wan Kenobi | Stewjon   | Human   |

# Tidy data, variables

## Structure

- **Columns hold variables**
- Rows
- Cells

| name           | homeworld | species |
|----------------|-----------|---------|
| Luke Skywalker | Tatooine  | Human   |
| R2-D2          | Naboo     | Droid   |
| Darth Vader    | Tatooine  | Human   |
| Obi-Wan Kenobi | Stewjon   | Human   |



# Tidy data, observations

## Structure

- Columns hold variables
- **Rows hold observations**
- Cells

| name           | homeworld | species |
|----------------|-----------|---------|
| Luke Skywalker | Tatooine  | Human   |
| R2-D2          | Naboo     | Droid   |
| Earth Vader    | Tatooine  | Human   |
| Obi Wan Kenobi | Stewjon   | Human   |

# Tidy data, values

## Structure

- Columns hold variables
- Rows hold observations
- **Cells hold values**

| name           | homeworld | species |
|----------------|-----------|---------|
| Luke Skywalker | Tatooine  | Human   |
| R2-D2          | Naboo     | Droid   |
| Darth Vader    | Tatooine  | Human   |
| Obi-Wan Kenobi | Stewjon   | Human   |

# dplyr recap

```
character_df
```

```
# A tibble: 4 x 3
  name          homeworld species
  <chr>         <chr>      <chr>
1 Luke Skywalker Tatooine   Human
2 R2-D2         Naboo      Droid
3 Darth Vader   Tatooine   Human
4 Obi-Wan Kenobi Stewjon    Human
```

# dplyr recap: select()

```
character_df %>%  
  select(name, homeworld)
```

```
# A tibble: 4 x 2  
  name          homeworld  
  <chr>         <chr>  
1 Luke Skywalker Tatooine  
2 R2-D2         Naboo  
3 Darth Vader   Tatooine  
4 Obi-Wan Kenobi Stewjon
```



# dplyr recap: filter()

```
character_df %>%  
  filter(homeworld == "Tatooine")
```

```
# A tibble: 2 x 3  
  name          homeworld species  
  <chr>         <chr>    <chr>  
1 Luke Skywalker Tatooine  Human  
2 Darth Vader   Tatooine  Human
```

# dplyr recap: mutate()

```
character_df %>%  
  mutate(is_human = species == "Human")
```

```
# A tibble: 4 x 4  
  name          homeworld species is_human  
  <chr>         <chr>    <chr>   <lgl>  
1 Luke Skywalker Tatooine   Human   TRUE  
2 R2-D2          Naboo     Droid   FALSE  
3 Darth Vader    Tatooine   Human   TRUE  
4 Obi-Wan Kenobi Stewjon   Human   TRUE
```


# dplyr recap: group\_by() and summarize()

```
character_df %>%  
  group_by(homeworld) %>%  
  summarize(n = n())
```

```
# A tibble: 3 x 2  
  homeworld      n  
  <chr>      <int>  
1 Naboo        1  
2 Stewjon      1  
3 Tatooine     2
```



<sup>1</sup> [magrittr.tidyverse.org](http://magrittr.tidyverse.org)



<sup>1</sup> [www.tidyverse.org](http://www.tidyverse.org)

# Multiple variables in a single column

```
population_df
```

```
# A tibble: 4 x 2
  country                population
  <chr>                  <dbl>
1 Brazil, South America 210.
2 Nepal, Asia           28.1
3 Senegal, Africa       15.8
4 Australia, Oceania    25.0
```

# Separating variables over two columns

```
population_df %>%  
  separate(country, into = c("country", "continent"), sep = ", ")
```

```
# A tibble: 4 x 3  
  country    continent    population  
  <chr>      <chr>          <dbl>  
1 Brazil    South America    210.  
2 Nepal     Asia             28.1  
3 Senegal   Africa           15.8  
4 Australia Oceania          25.0
```

# Let's practice!

RESHAPING DATA WITH TIDYR



# Columns with multiple values

RESHAPING DATA WITH TIDYR



**Jeroen Boeye**

Head of Machine Learning, Faktion

# Two variables in a single column

```
netflix_df
```

```
# A tibble: 637 x 3
  title                type    duration
  <chr>                <chr>    <chr>
1 Article 15          Movie    125 min
2 Kill Me If You Dare Movie    100 min
3 The Spy             TV Show  1 Seasons
4 The World We Make   Movie    108 min
5 Watchman            Movie     93 min
```

# Converting separated columns' data types

```
netflix_df %>%  
  separate(duration, into = c("value", "unit"), convert = TRUE)
```

```
# A tibble: 5 x 4  
  title          type    value unit  
  <chr>         <chr>  <int> <chr>  
1 Article 15     Movie   125 min  
2 Kill Me If You Dare Movie   100 min  
3 The Spy       TV Show    1 Seasons  
4 The World We Make Movie   108 min  
5 Watchman      Movie    93 min
```

# dplyr aggregation recap

```
netflix_df %>%  
  separate(duration, into = c("value", "unit"), convert = TRUE) %>%  
  group_by(type, unit) %>%  
  summarize(mean_duration = mean(value))
```

```
# A tibble: 2 x 3  
# Groups:   type [2]  
  type      unit mean_duration  
  <chr>   <chr>         <dbl>  
1 Movie    min           98.6  
2 TV Show Seasons       1.85
```

# Separating variables over columns

| title | type | duration |
|-------|------|----------|
|       |      |          |
|       |      |          |

| title | type | value | unit |
|-------|------|-------|------|
|       |      |       |      |
|       |      |       |      |

# Combining multiple columns into one

```
star_wars_df
```

```
# A tibble: 4 x 2
  given_name family_name
  <chr>      <chr>
1 Luke      Skywalker
2 Han       Solo
3 Leia      Organa
4 R2        D2
```

# Combining multiple columns into one

```
star_wars_df %>%  
  unite("name", given_name, family_name)
```

```
# A tibble: 4 x 1  
  name  
  <chr>  
1 Luke_Skywalker  
2 Han_Solo  
3 Leia_Organa  
4 R2_D2
```

# Combining multiple columns into one

```
star_wars_df %>%  
  unite("name", given_name, family_name, sep = " ")
```

```
# A tibble: 4 x 1  
  name  
  <chr>  
1 Luke Skywalker  
2 Han Solo  
3 Leia Organa  
4 R2 D2
```



# Multiple values in a single cell

```
drink_df
```

```
# A tibble: 2 x 2  
  drink      ingredients  
  <chr>      <chr>  
1 Chocolate milk milk, chocolate, sugar  
2 Orange juice oranges, sugar
```

# Multiple values in a single cell

## Netflix data

| title | type | duration |
|-------|------|----------|
|       |      |          |
|       |      |          |

## Drinks data

| drink | ingredients |   |   |
|-------|-------------|---|---|
| A     | 1           | 2 | 3 |
| B     | 1           | 2 |   |

# Multiple values in a single cell

## Netflix data

| title | type | duration |
|-------|------|----------|
|       |      |          |
|       |      |          |

## Values to variables

| title | type | value | unit |
|-------|------|-------|------|
|       |      |       |      |
|       |      |       |      |

## Drinks data

| drink | ingredients |   |   |
|-------|-------------|---|---|
| A     | 1           | 2 | 3 |
| B     | 1           | 2 |   |

# Multiple values in a single cell

## Netflix data

| title | type | duration |
|-------|------|----------|
|       |      |          |
|       |      |          |

## Values to variables

| title | type | value | unit |
|-------|------|-------|------|
|       |      |       |      |
|       |      |       |      |

## Drinks data

| drink | ingredients |   |   |
|-------|-------------|---|---|
| A     | 1           | 2 | 3 |
| B     | 1           | 2 |   |

## Values to observations

| drink | ingredients |
|-------|-------------|
| A     | 1           |
| A     | 2           |
| A     | 3           |
| B     | 1           |
| B     | 2           |

# Separating values over rows

```
drink_df %>%  
  separate_rows(ingredients, sep = ", ")
```

```
# A tibble: 5 x 2  
  drink      ingredients  
  <chr>      <chr>  
1 Chocolate milk milk  
2 Chocolate milk chocolate  
3 Chocolate milk sugar  
4 Orange juice oranges  
5 Orange juice sugar
```

# Counting ingredients

```
drink_df %>%  
  separate_rows(ingredients, sep = ", ") %>%  
  count(drink)
```


```
# A tibble: 2 x 2  
  drink      n  
  <chr>    <int>  
1 Chocolate milk    3  
2 Orange juice     2
```

```
drink_df %>%  
  separate_rows(ingredients, sep = ", ") %>%  
  count(ingredients)
```

```
# A tibble: 4 x 2  
  ingredients      n  
  <chr>          <int>  
1 chocolate      1  
2 milk           1  
3 oranges        1  
4 sugar          2
```

# Visualizing ingredients

```
drink_df %>%  
  separate_rows(ingredients, sep = ", ") %>%  
  ggplot(aes(x=drink, fill=ingredients)) +  
  geom_bar()
```



# Let's practice!

RESHAPING DATA WITH TIDYR



# Missing values

RESHAPING DATA WITH TIDYR



**Jeroen Boeye**

Head of Machine Learning, Faktion

# Missing values in R

## NA = Not Available

```
# A tibble: 5 x 4
  drink      ingredient quantity unit
  <chr>      <chr>          <int> <chr>
1 Chocolate milk milk           1 L
2 Chocolate milk chocolate     100 g
3 Chocolate milk sugar         20 g
4 Orange juice oranges          3 NA
5 Orange juice sugar          20 g
```

# Imputing with a default value: `replace_na()`

```
moon_df
```

```
# A tibble: 4 x 2
  year people_on_moon
<int>      <int>
1  1969          4
2  1970         NA
3  1971          4
4  1972          4
5  1973         NA
```

# Imputing with a default value: `replace_na()`

```
moon_df %>%  
  replace_na(list(people_on_moon = 0L))
```

```
# A tibble: 4 x 2  
  year people_on_moon  
  <int>         <int>  
1  1969             4  
2  1970             0  
3  1971             4  
4  1972             4  
5  1973             0
```

```
typeof(0L)
```

```
[1] "integer"
```

```
typeof(0)
```

```
[1] "double"
```

# Imputing with the most recent value: fill()

```
cumul_moon_df
```

```
# A tibble: 5 x 3
  year people_on_moon total_people_on_moon
  <int>         <int>         <int>
1  1969             4             4
2  1970            NA            NA
3  1971             4             8
4  1972             4            12
5  1973            NA            NA
```

# Imputing with the most recent value: fill()

```
cumul_moon_df %>%  
  fill(total_people_on_moon)
```

```
# A tibble: 5 x 3  
  year people_on_moon total_people_on_moon  
  <int>         <int>         <int>  
1  1969             4             4  
2  1970            NA             4  
3  1971             4             8  
4  1972             4            12  
5  1973            NA            12
```

# fill() imputation options

```
cumul_moon_df %>%  
  fill(total_people_on_moon, .direction = "down")
```

```
# A tibble: 5 x 3  
  year people_on_moon total_people_on_moon  
  <int>         <int>         <int>  
1  1969             4             4  
2  1970            NA             4  
3  1971             4             8  
4  1972             4            12  
5  1973            NA            12
```

# fill() imputation options

```
cumul_moon_df %>%  
  fill(total_people_on_moon, .direction = "up")
```

```
# A tibble: 5 x 3  
  year people_on_moon total_people_on_moon  
  <int>         <int>         <int>  
1  1969             4             4  
2  1970            NA             8  
3  1971             4             8  
4  1972             4            12  
5  1973            NA            NA
```



# Removing rows with missing values: drop\_na()

```
moon_df %>%  
  drop_na()
```

```
# A tibble: 3 x 2  
  year people_on_moon  
  <int>         <int>  
1  1969             4  
2  1971             4  
3  1972             4
```

# drop\_na() caveats

```
mars_df
```

```
# A tibble: 5 x 3
  year people_on_moon people_on_mars
<int>         <int> <int>
1  1969             4 NA
2  1970            NA NA
3  1971             4 NA
4  1972             4 NA
5  1973            NA NA
```

# drop\_na() caveats

```
mars_df %>%  
  drop_na()
```

```
# A tibble: 0 x 3  
# ... with 3 variables: year <int>, people_on_moon <int>, people_on_mars <int>
```

# drop\_na() caveats

```
mars_df %>%  
  drop_na(people_on_moon)
```

```
# A tibble: 3 x 3  
  year people_on_moon people_on_mars  
  <int>         <int> <int>  
1  1969             4 NA  
2  1971             4 NA  
3  1972             4 NA
```

# Let's practice!

RESHAPING DATA WITH TIDYR